

Customer Acquisition Case Study

Cynthia Farias, Maria Morinigo, Robert Perez, Rubina Saya

4/23/2021

Executive Summary

In this case study, the aim is to fit several predictive models to find the best model that accurately predicts the acquisition of new customers. Our analysis resulted in the `random forest` model providing the highest accuracy rate of 82.29%. This accuracy was compared to logistic and decision tree models, which resulted in an accuracy of 80% and 71.33%, respectively. Even though the `logistic model`'s accuracy is comparable, the `random forest` was selected as the final model due to its ability to deal with multicollinearity issues, non-linear relationships, missing values, outliers, and unbalanced data.

Further, it was noted that the most significant variable identified using the three models consistently included `industry`, `revenue`, and `employees`, with `random forest` also identifying `acq_exp` as having significance.

On customer observations predicted as `acquired` using the `random forest` model, another tuned `random forest` was utilized to identify the most significant variables for customer duration. The model resulted in the `ret_exp_sq`, `ret_exp`, `freq_sq`, and `freq` variables being the most significant in determining the duration of the customer with the company.

Overall, the `random forest` proved to be a reliable model in predicting customer acquisition. The resulting information from this case study related to customer acquisition and duration will provide companies with valuable insights to establish informed targeted marketing strategies to gain business, reduce cost, and increase profitability.

Problem

For a company to be successful, it must understand which variables impact customer acquisition and retention. The purpose of this case study is to identify the best model to predict customer acquisition and then to apply the selected model to classify the most significant explanatory variables impacting both customer acquisition and duration. This will allow the company to take appropriate actions to ensure that customers are acquired and then invest in resources to increase the customers' tenure with the company. The company may use the knowledge gained to deploy targeted marketing strategies, including discounts and other incentives, to gain and retain customers.

The goal of this case study is to answer the following questions:

- Which predictive model between Random Forest, Logistic, and Decision Tree can predict customer acquisition Retention data set to predict which customers will be acquired?
- Which variables are significant in determining new customer acquisition and duration of the customers with the company?
- If a customer is acquired, for how long is that customer retained?

The remaining sections of this report will look at the existing literature related to models used to predict customer retention and churn, discuss the data exploration and preprocessing, methodologies applied to develop the best model, and present the results and recommendations.


Review of Related Literature

While still a relatively modern machine learning technique, random forests have seen extensive use since first introduced by Breiman in his paper Random Forests (2001). Specifically, `random forests` have been applied to the problem of predicting customer retention by a handful of researchers. Larivière and Poel compared a model based on `random forests` to traditional prediction models and found it superior to both `logistic` and `linear regression models` when measured by predictive accuracy (2005). More recently, Sabbeh compared a wide range of popular techniques for customer "churn" prediction in a "battle royal" that tested logistic regression, decision tree, Naïve Bayesian, Support Vector Machine (SVM), K-Nearest Neighbor (KNN), ensemble learning, which includes `random forest`, artificial neural network (ANN), and linear discriminant analysis (LDA). His results indicated that random forest was superior to other techniques with an accuracy of roughly 96%. He did note that SVM was close behind with an accuracy of 94%.

Schaeffer and Sanchez found that the performance of a `random forest` model was superior to an `SVM` model but determined that the accuracy gain was small enough to make SVM a viable option as its performance was much better by several factors (2020). While this difference amounted to only a few seconds in running the model on their dataset, they noted that the additional accuracy of `random forests` may not be worth the cost in time for large datasets. For all other situations where accuracy results in increased profit and time is not a significant factor, the `random forests` technique appears to be emerging as the go-to model for customer retention prediction.

Methods

The case study aims to accurately predict customer retention and acquisition to target the right customers and decrease the cost of marketing campaigns. To fit the models, we are using decision trees, logistic regression, and random forest. We use the `acquisitionRetention` data set to predict which customers will be acquired. Prior to fitting the models, we performed some data cleaning and looked for missing values and multicollinearity. Additionally, we converted `crossbuy`, `industry`, and `acquisition` into factors. The correlation plot below shows some multicollinearity between independent variables. For instance `duration` has high correlation with `profit` and `ret_exp`, `profit` has a high correlation with `ret_exp`, and `acq_exp` has a high correlation with `ac_exp_sq`.



Next, we split the data into a 70 percent training set and 30 percent testing set. We fit our first model using a decision tree. This model technique is a supervised learning algorithm that can be used in regression problems. The response variable is `acquisition`, and the independent variables were `acq_exp`, `industry`, `revenue`, and `employees`. Using decision trees has some advantages. One of the advantages of a Decision Tree is that the output is easy to interpret.

Decision Tree is also helpful in data exploration because it identifies the most significant variables. For instance, the decision tree model indicated that `employees`, `acq_exp`, and `revenue` were more important. Another advantage is that it requires less data cleaning, and outliers have no influence on the model, and it can handle both numerical and categorical variables. In this case study, the response variable, `acquisition`, is categorical. However, decision trees have some disadvantages as well. Overfitting can be a problem when using decision trees. Some techniques can be performed to improve the performance of decision trees like bagging, random forest, and boosting (Arka Roy, DA 6813. DATA ANALYTICS APPLICATIONS SPRING 2021. PowerPoint).

The next technique used was logistic regression. We fit our `glm` model using the following predictors: `acq_exp`, `acq_exp_sq`, `industry`, `revenue`, and `employees`. In the original data, the `acquisition` variable was numerical, with 1 if the prospect was acquired and 0 otherwise. The response variable was converted to a factor to fit a logistic model. Next, we use the `vif()` function to calculate the variance inflation factor to measure how much the variance of a regression coefficient is inflated due to multicollinearity. We eliminated `acq_exp`, which had a $vif > 5$. That left us with a model using the following predictors: `acq_exp_sq`, `industry`, `revenue`, and `employees`. The model's performance increased from 0.7133 accuracy in the decision tree model to 0.8 accuracy on the logistic model. Although the logistic regression model outperformed the decision tree model and identified the variables that were significant to the model, logistic regression models have some limitations. For instance, logistic regression cannot show complex relationships, and multicollinearity can cause issues between independent variables.

The last technique used in this case study was random forest. Random forest is a machine learning method used in regression and classification tasks. Random forest is good at handling missing values, outliers, unbalanced data, and other steps in the data exploration process. Similar to decision trees, random forest is also good at identifying important variables. However, one main disadvantage of random forest is overfitting. To grow our forest, first we decided to fit our forest model by manually selecting the predictors. Our first forest model had the following predictors: `acq_exp`, `acq_exp_sq`, `industry`, and `revenue`. This model gave us an overall error rate of 21.43%.

The next forest model had the following predictors: `acq_exp_sq`, `industry`, and `revenue`. This model gave us an overall error rate of 18.29%. A third forest model was fit with the following predictors: `acq_exp`, `industry`, and `revenue`. And this model gave us an overall error rate of 18.29%, similar to the second model. The next step was to optimize our selection process. To identify the optimal set of hyper-parameters based on `oob_error`, we established a list of possible values for the hyper-parameters, we created a data frame containing all combinations, we created an empty vector to store `oob_error` values, and wrote a for loop. We fitted a new forest using the optimal hyper-parameters with a resulting overall error rate of 17.71%.

Finally, we used the predictions from the random forest hyper-parameters model to create a subset of data where customer were predicted to be acquired. We split this subset data set into training and testing sets and fit a model using random forest. Our response variable for this model was `duration`. We followed the same process as before, first, we made a manual selection of independent variables, and then we optimized our selection, in the same way, explained earlier. The model with the optimized parameters gave us better results (which is presented in the result section).

Data

The Customer Acquisition dataset is a small portion of 500 customers from a typical B2B firm, where all the customers made their first purchase with the firm at the same time (Kumar & Peterson, 2012). The data is embedded in the `SMCRM` package in R. For this case analysis, we created subsets of the database: a training dataset and a testing dataset. The training dataset contains 350 observations and 14 predicting variables. The testing dataset contains 150 observations and 14 variables. The test dataset is a holdout sample representative of the entire population, and it is used to assess the performance of the model. The variables are as follows:

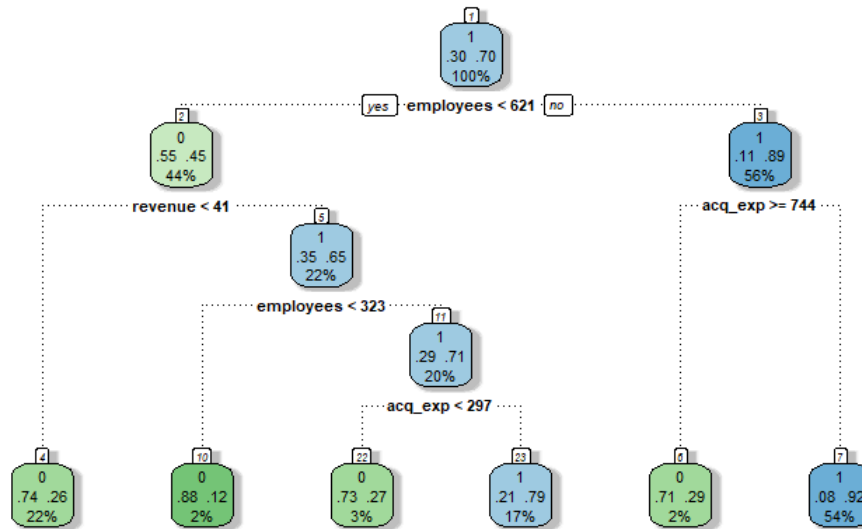
	Variable Name	Description
1.	<code>customer</code>	customers number (from 1 to 500)
2.	<code>acquisition</code>	1 if the prospect was acquired, 0 otherwise
3.	<code>duration</code>	number of days the customer was a customer of the firm, 0 if <code>acquisition == 0</code>
4.	<code>profit</code>	customer lifetime value (CLV) of a given customer, $-(Acq_Exp)$ if the customer is not acquired
5.	<code>acq_exp</code>	total dollars spent on trying to acquire this prospect
6.	<code>ret_exp</code>	total dollars spent on trying to retain this customer
7.	<code>acq_exp_sq</code>	square of the total dollars spent on trying to acquire this prospect
8.	<code>ret_exp_sq</code>	square of the total dollars spent on trying to retain this customer
9.	<code>freq</code>	number of purchases the customer made during that customer's lifetime with the firm, 0 if <code>acquisition == 0</code>
10.	<code>freq_sq</code>	square of the number of purchases the customer made during that customer's lifetime with the firm
11.	<code>crossbuy</code>	number of product categories the customer purchased from during that customer's lifetime with the firm, 0 if <code>acquisition = 0</code>
12.	<code>sow</code>	Share-of-Wallet; percentage of purchases the customer makes from the given firm given the total amount of purchases across all firms in that category
13.	<code>industry</code>	1 if the customer is in the B2B industry, 0 otherwise
14.	<code>revenue</code>	annual sales revenue of the prospect's firm (in millions of dollar)
15.	<code>employees</code>	number of employees in the prospect's firm

Prior to training the model based on the dataset, some data cleaning had to be done. The data preprocessing for this case study entailed looking for missing values, looking at multicollinearity, and possible elimination of irrelevant variables. This dataset was relatively clean and did not have any missing values. To better understand and make inferences about the data, some numerical and graphical summaries were performed. The Correlation Matrix and Correlation Plot were utilized to assess any variables that had a relationship with each other. It revealed that `profit`, and `ret_exp`, `acq_exp`, and `acq_exp_sq`, `ret_exp`, and `ret_exp_sq`, `freq`, and `freq_sq`, had a strong correlation.

Additionally, while creating our models, we did not include some variables that were not relevant for our analysis. We used `as.factor()` to convert `industry`, `crossbuy`, and `Acquisition` into factor variables for Decision Tree, Logistic Regression, and Random Forest Models.

Results

For the initial model selection of acquisition, when constructing the tree model, the only variables included were `acq_exp`, `industry`, `revenue` and `employees`. Other variables were excluded due to providing perfect separation for the acquisition variable. For example, the `duration` variable is equal to zero for all customers that were not acquired. As a result, the `duration` variable does not allow for deeper insight of other variables' effects, as we would only have to determine if a customer's duration is greater than zero to know if they were acquired or not.

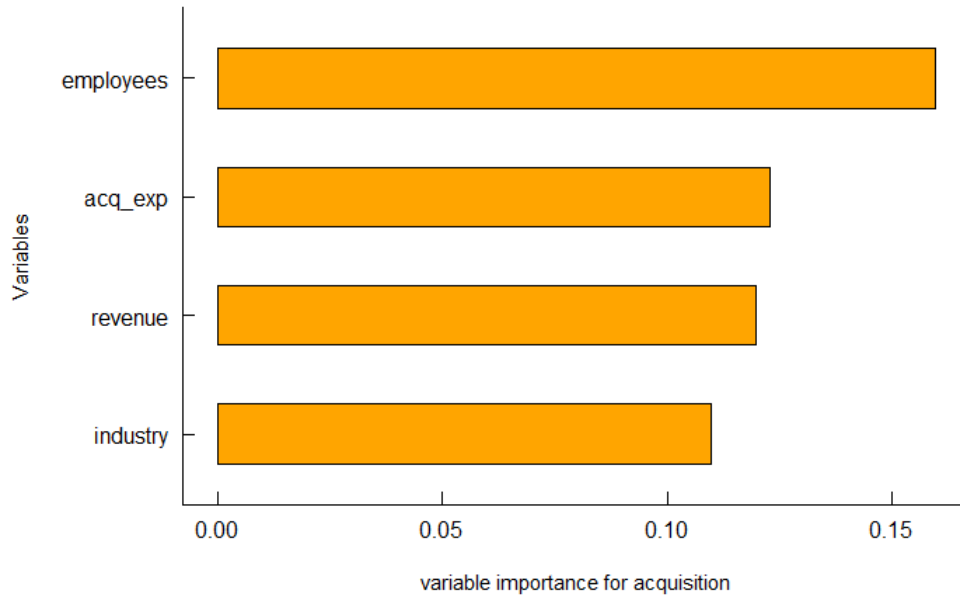


The resulting tree model used an initial split at `employees < 621`, with further splits at `revenue`, `acq_exp`, and `employees`, once more. The accuracy of the model was roughly 71.33%, with 73.58% sensitivity and 65.91% specificity.


A similar approach was used when constructing a logistic model for the data. Initially the variables used were the same as the tree model, with the addition of the `acq_exp_sq`. After running the model, a high correlation was found between the `acq_exp` and `acq_exp_sq`, the `acq_exp` was removed from the model and the model was rerun. The resulting model showed that the variables `industry`, `revenue`, and `employees` all had significant effects on customer acquisition. The model displayed an overall accuracy of 80%, with sensitivity of 77.88% and specificity of 86.49%.

For the forest model, two different approaches were taken. The first approach was a manual selection, while the second approach used an optimized selection to determine which hyper parameters were appropriate for the model.

The manual model created used the same variables as the initial tree model as an input formula. The resulting metrics showed an accuracy of roughly 81.71% with a sensitivity of 91.02% and a 59.05% specificity. This model showed a large improvement over the previous two models created. Additionally, the `employees` variable showed the most importance of all variables, with roughly 0.0597 importance overall. `acq_exp` and `revenue` showed similar importance, 0.0229 and 0.0199 respectively. Finally, the `industry` variable had the least amount of importance for acquisition of a customer. The graph below shows these metrics with a constant of 0.1 added for easier viewing.




The optimized model was run with `mtry`, the number of variables randomly selected for splitting a node, equal to 1, a minimum terminal node size of 1, and 1000 trees. The results show an accuracy of 82.29%, 55.24% specificity, and a 99.94% sensitivity, a further improvement from the manually selected random forest model. The importance values for `employees`, `acq_exp_sq`, `revenue`, and `industry` were 0.0644, 0.0271, 0.0265, and 0.0135, respectively. The graphic again shows the values with a constant of 0.1 added for ease of viewing. Although the model showed average results for properly predicting negative cases, it displayed the best overall accuracy, as well as having excellent ability to predict positive cases.



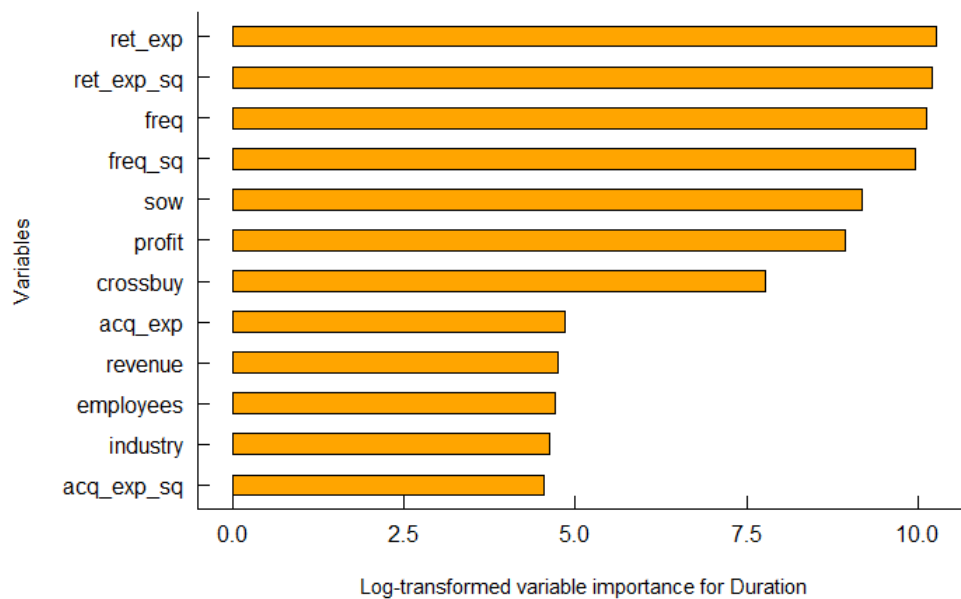
Predictions were then carried out using this model on the complete dataset. From these predictions, data points with a positive prediction of acquisition were separated from the dataset. The average `duration` for this subsetted data was roughly 949.32 days. A random forest model with manual selection and model with optimized hyper parameters was then created with `duration` as the response variable, including all variables except for acquisition. The model with manual selection produced an error rate of 1404.75. The variables with the greatest importance are `ret_exp_sq`, `ret_exp`, `freq_sq`, and `freq`.

```
{r}
summary(Acquired.df$duration)
```



Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.0	869.2	1038.0	949.3	1207.0	1673.0





The same variables were used as the input formulas to optimize the hyper parameters for the second random forest model. The resulting optimized hyper parameters were an `mtry` of 6, minimum terminal node size of 1, and 1000 trees. The optimized model returned an error rate of 1033.72, with a mean square error of 1525.854 and a mean absolute error of 25.368. The variables with greatest importance were the `ret_exp_sq`, `ret_exp`, `freq_sq`, and `freq` variables. Interactions between these four variables were inspected, and the model rerun with these interaction variables. However, the new rerun model generated the same optimized hyperparameters and the same results as the model that did not include interactions.



Partial dependence plots were then created for each of the variables with the greatest importance. From these plots, it can be observed that there are no linear relations between the variables and the duration of an acquired customer.





Conclusion & Recommendations

In conclusion, a tuned Random Forest model outperformed both logistic and decision tree models in predicting new customer acquisition. The variables that significantly impact customer acquisition include `employees`, `acq_exp_sq`, `revenue`, and `industry`. It is possible that more and more customers are interested in a better quality of service, it makes sense that the number of employees has the most significant effect on customer acquisition, followed by the dollars invested in attaining the customer.

Based on our final results, we identified that the `ret_exp_sq`, `ret_exp`, `freq_sq`, and `freq` variables significantly impact the customer `duration` variable. Additionally, we determined that the average duration of the customer is around 949.32 days, i.e., 2.6 years.

The company would gain insightful information if, in addition to creating predictive models to determine the significant variables impacting customer acquisition and retention, it also analyzed customer demographics to understand the customer base better. This knowledge will further allow the company to deploy more efficient and targeted approaches to utilize scarce resources to gain business while reducing the cost of marketing campaigns and increasing profitability.

References

- Breiman, L. (2001). Random Forests. *Machine Learning*, 45(1), 5–32
- Kumar, V., & Peterson, J. (2012). Statistical Methods in Customer Relationship Management. Wiley.
- Larivière, B., & Van den Poel, D. (2005). Predicting customer retention and profitability by using random forests and regression forests techniques. *Expert Systems with Applications*, 29(2), 472–484.
- Schaeffer, S., & Rodriguez Sanchez, S. (2020). Forecasting client retention — A machine-learning approach. *Journal of Retailing and Consumer Services*, 52, 101918–.
- Sabbeh, S. (2018). Machine-Learning Techniques for Customer Retention: A Comparative Study. *International Journal of Advanced Computer Science and Applications*, 9(2), 273-281.

Data Exploration

Data Importing

```
data("acquisitionRetention")
```

```
View(acquisitionRetention)  
str(acquisitionRetention)
```

```
## 'data.frame': 500 obs. of 15 variables:  
## $ customer : num 1 2 3 4 5 6 7 8 9 10 ...  
## $ acquisition: num 1 1 1 0 1 1 1 1 0 0 ...  
## $ duration : num 1635 1039 1288 0 1631 ...  
## $ profit : num 6134 3524 4081 -638 5446 ...  
## $ acq_exp : num 694 460 249 638 589 ...  
## $ ret_exp : num 972 450 805 0 920 ...  
## $ acq_exp_sq : num 480998 211628 62016 407644 346897 ...  
## $ ret_exp_sq : num 943929 202077 648089 0 846106 ...  
## $ freq : num 6 11 21 0 2 7 15 13 0 0 ...  
## $ freq_sq : num 36 121 441 0 4 49 225 169 0 0 ...  
## $ crossbuy : num 5 6 6 0 9 4 5 5 0 0 ...  
## $ sow : num 95 22 90 0 80 48 51 23 0 0 ...  
## $ industry : num 1 0 0 0 0 1 0 1 0 1 ...  
## $ revenue : num 47.2 45.1 29.1 40.6 48.7 ...  
## $ employees : num 898 686 1423 181 631 ...
```

Data Cleaning

```
sapply(acquisitionRetention, function(x) sum(is.na(x)))
```

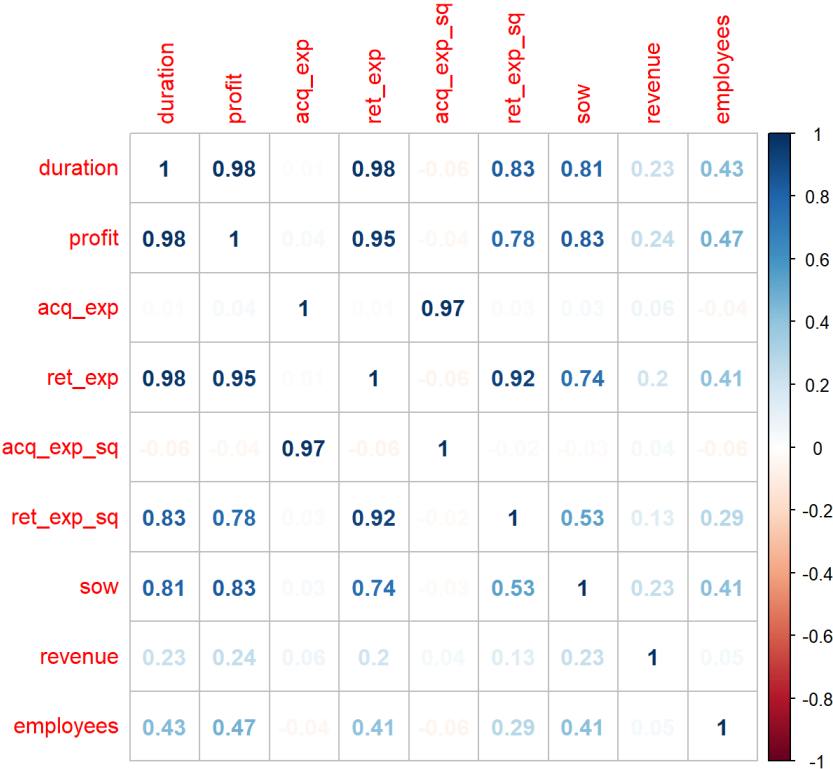
```
## customer acquisition duration profit acq_exp ret_exp  
## 0 0 0 0 0 0  
## acq_exp_sq ret_exp_sq freq freq_sq crossbuy sow  
## 0 0 0 0 0 0  
## industry revenue employees  
## 0 0 0
```

```
acquisitionRetention <- acquisitionRetention[,2:15]  
acquisitionRetention$crossbuy <- as.factor(acquisitionRetention$crossbuy)  
acquisitionRetention$industry <- as.factor(acquisitionRetention$industry)  
acquisitionRetention$acquisition <- as.factor(acquisitionRetention$acquisition)  
str(acquisitionRetention)
```

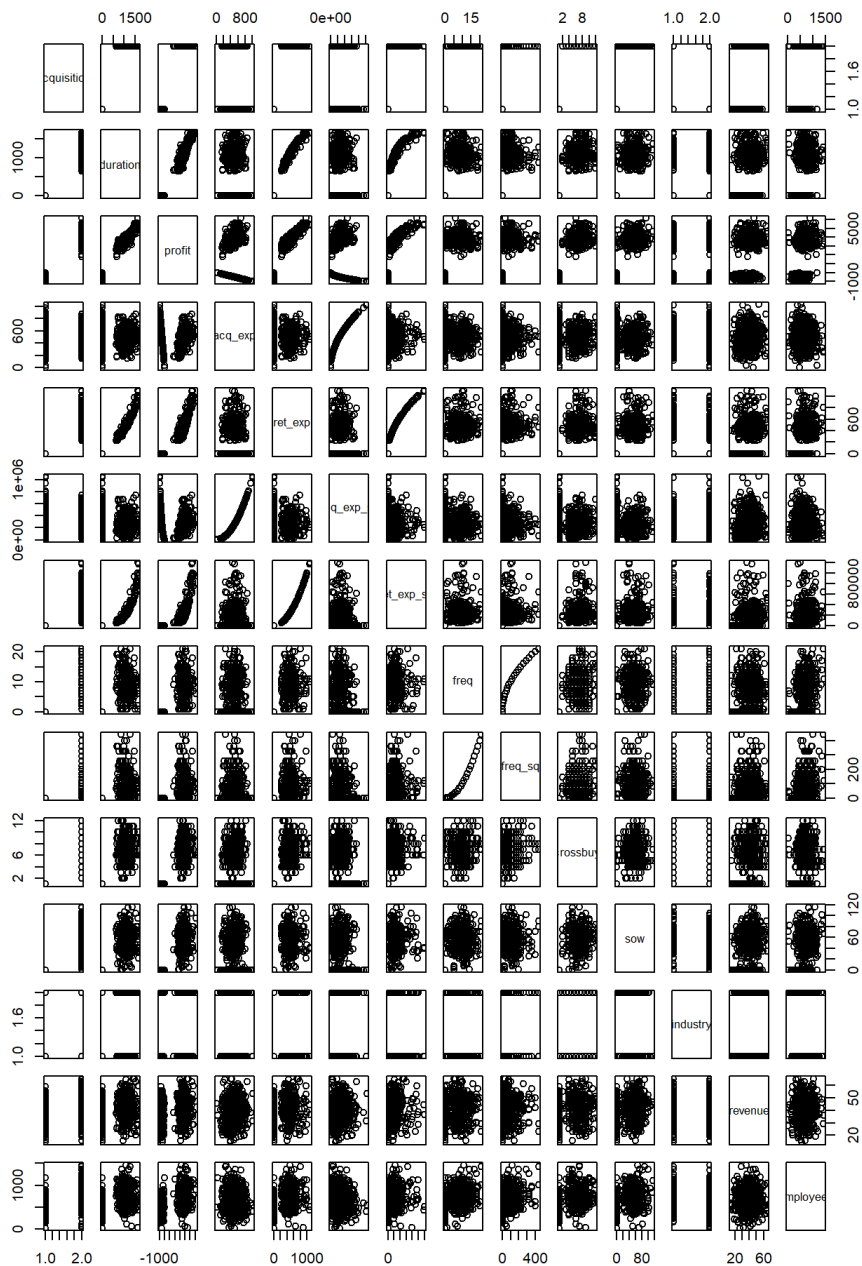
```
## 'data.frame': 500 obs. of 14 variables:  
## $ acquisition: Factor w/ 2 levels "0","1": 2 2 2 1 2 2 2 2 1 1 ...  
## $ duration : num 1635 1039 1288 0 1631 ...  
## $ profit : num 6134 3524 4081 -638 5446 ...  
## $ acq_exp : num 694 460 249 638 589 ...  
## $ ret_exp : num 972 450 805 0 920 ...  
## $ acq_exp_sq : num 480998 211628 62016 407644 346897 ...  
## $ ret_exp_sq : num 943929 202077 648089 0 846106 ...  
## $ freq : num 6 11 21 0 2 7 15 13 0 0 ...  
## $ freq_sq : num 36 121 441 0 4 49 225 169 0 0 ...  
## $ crossbuy : Factor w/ 12 levels "0","1","2","3",...: 6 7 7 1 10 5 6 6 1 1 ...  
## $ sow : num 95 22 90 0 80 48 51 23 0 0 ...  
## $ industry : Factor w/ 2 levels "0","1": 2 1 1 1 1 2 1 2 1 2 ...  
## $ revenue : num 47.2 45.1 29.1 40.6 48.7 ...  
## $ employees : num 898 686 1423 181 631 ...
```

Data Correlations

```
corrplot(cor(acquisitionRetention[,c(2:7,11,13:14)]), method = "number")
```



```
pairs(acquisitionRetention)
```



Data Splitting

Split for Acquisition prediction

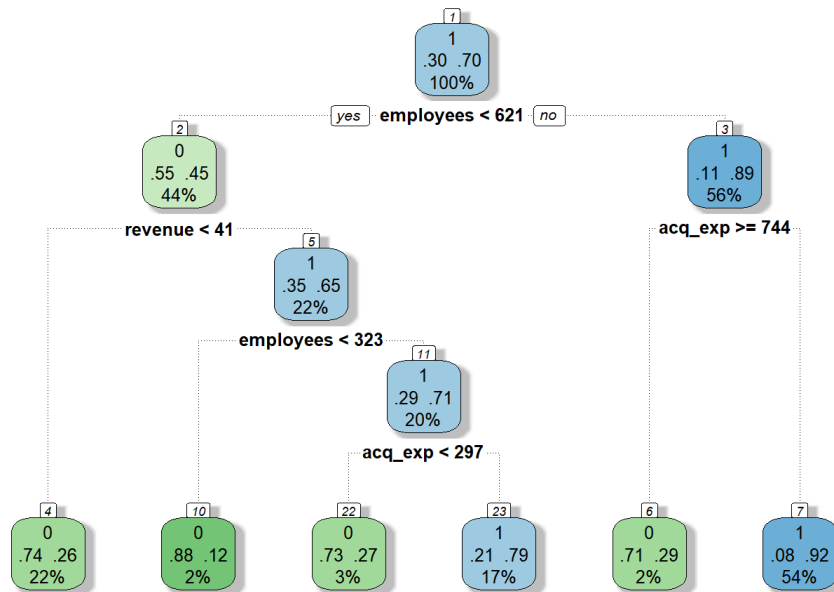
```
set.seed(123)
idx.train <- sample(1:nrow(acquisitionRetention), size = 0.7 * nrow(acquisitionRetention))
train.df <- acquisitionRetention[idx.train,]
test.df <- acquisitionRetention[-idx.train,]
```

Model Creation & Predictions

Tree Model - Acquisition

```
set.seed(123)
dt.model <- rpart(acquisition ~ acq_exp + industry + revenue + employees, data = train.df) # simple DT model

rattle::fancyRpartPlot(dt.model, sub = "") # visualize the DT
```



```
predicted.acquisition <- predict(dt.model, newdata = test.df, type = "class")
View(predicted.acquisition)
```

```
caret::confusionMatrix(as.factor(test.df$acquisition), as.factor(predicted.acquisition), positive='1')
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 29 28
##           1 15 78
##
##           Accuracy : 0.7133
##           95% CI : (0.6339, 0.7841)
##           No Information Rate : 0.7067
##           P-Value [Acc > NIR] : 0.46919
##
##           Kappa : 0.3635
##
## Mcnemar's Test P-Value : 0.06725
##
##           Sensitivity : 0.7358
##           Specificity : 0.6591
##           Pos Pred Value : 0.8387
##           Neg Pred Value : 0.5088
##           Prevalence : 0.7067
##           Detection Rate : 0.5200
##           Detection Prevalence : 0.6200
##           Balanced Accuracy : 0.6975
##
##           'Positive' Class : 1
##
```

Logistic Model - Acquisition

```
set.seed(123)
glm.model <- glm(acquisition ~ acq_exp + acq_exp_sq + industry + revenue + employees, data = train.df, family = "binomial")
summary(glm.model)
```

```
##
## Call:
## glm(formula = acquisition ~ acq_exp + acq_exp_sq + industry +
##     revenue + employees, family = "binomial", data = train.df)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2237  -0.2916   0.1958   0.5023   2.4956
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.524e+01  1.879e+00  -8.108 5.15e-16 ***
## acq_exp      3.286e-02  5.336e-03   6.158 7.39e-10 ***
## acq_exp_sq   -3.267e-05  5.334e-06  -6.124 9.15e-10 ***
## industry1    1.695e+00  3.585e-01   4.728 2.27e-06 ***
## revenue      8.734e-02  1.784e-02   4.895 9.84e-07 ***
## employees    7.480e-03  1.011e-03   7.399 1.37e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 427.61  on 349  degrees of freedom
## Residual deviance: 231.92  on 344  degrees of freedom
## AIC: 243.92
##
## Number of Fisher Scoring iterations: 6
```

```
car::vif(glm.model)
```

```
##   acq_exp acq_exp_sq  industry  revenue  employees
##  28.382762 28.030374  1.129658  1.082911  1.265077
```

```
set.seed(123)
glm.model2 <- glm(acquisition ~ acq_exp_sq + industry + revenue + employees, data = train.df, family = "binomial")
summary(glm.model2)
```

```
##
## Call:
## glm(formula = acquisition ~ acq_exp_sq + industry + revenue +
##     employees, family = "binomial", data = train.df)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.1886  -0.5368   0.3057   0.6448   2.4065
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -7.097e+00  9.819e-01 -7.228 4.91e-13 ***
## acq_exp_sq   -6.132e-07  8.590e-07 -0.714   0.475
## industry1    1.303e+00  3.092e-01  4.215 2.50e-05 ***
## revenue      8.807e-02  1.623e-02  5.425 5.80e-08 ***
## employees    6.561e-03  8.679e-04  7.559 4.06e-14 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 427.61  on 349  degrees of freedom
## Residual deviance: 285.96  on 345  degrees of freedom
## AIC: 295.96
##
## Number of Fisher Scoring iterations: 5
```

```
car::vif(glm.model2)
```

```
## acq_exp_sq  industry  revenue  employees
##   1.018414   1.087649   1.089387   1.120444
```

```
glm.preds <- predict(glm.model2, newdata = test.df, type = "response")
test.df$PredChoice = ifelse(glm.preds >= 0.5, 1, 0)
test.df$PredChoice = as.factor(test.df$PredChoice)
```

```
caret::confusionMatrix(as.factor(test.df$acquisition), as.factor(test.df$PredChoice), positive = '1')
```



```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 32 25
##           1   5 88
##
##           Accuracy : 0.8
##           95% CI : (0.727, 0.8608)
##           No Information Rate : 0.7533
##           P-Value [Acc > NIR] : 0.1073527
##
##           Kappa : 0.5446
##
## Mcnemar's Test P-Value : 0.0005226
##
##           Sensitivity : 0.7788
##           Specificity : 0.8649
##           Pos Pred Value : 0.9462
##           Neg Pred Value : 0.5614
##           Prevalence : 0.7533
##           Detection Rate : 0.5867
##           Detection Prevalence : 0.6200
##           Balanced Accuracy : 0.8218
##
##           'Positive' Class : 1
##
```

Forest Model - Acquisition

theme for nice plotting

```
theme_nice <- theme_classic()+
  theme(
    axis.line.y.left = element_line(colour = "black"),
    axis.line.y.right = element_line(colour = "black"),
    axis.line.x.bottom = element_line(colour = "black"),
    axis.line.x.top = element_line(colour = "black"),
    axis.text.y = element_text(colour = "black", size = 12),
    axis.text.x = element_text(color = "black", size = 12),
    axis.ticks = element_line(color = "black")) +
  theme(
    axis.ticks.length = unit(-0.25, "cm"),
    axis.text.x = element_text(margin=unit(c(0.5,0.5,0.5,0.5), "cm")),
    axis.text.y = element_text(margin=unit(c(0.5,0.5,0.5,0.5), "cm")))
```

Manual Selection

```
set.seed(123)
forest1 <- rfsrc(acquisition ~ acq_exp + acq_exp_sq + industry + revenue + employees,
  data = train.df,
  importance = TRUE,
  ntree = 1000)

forest1
```

```
##           Sample size: 350
##       Frequency of class labels: 105, 245
##           Number of trees: 1000
##       Forest terminal node size: 1
##       Average no. of terminal nodes: 49.886
## No. of variables tried at each split: 3
##       Total no. of variables: 5
##       Resampling used to grow trees: swor
##       Resample size used to grow trees: 221
##           Analysis: RF-C
##           Family: class
##       Splitting rule: gini *random*
##       Number of random split points: 10
##       Normalized brier score: 56.97
##           AUC: 85.2
##       Error rate: 0.21, 0.43, 0.12
##
## Confusion matrix:
##
##       predicted
## observed  0   1 class.error
##       0 60  45      0.4286
##       1 30 215      0.1224
##
## Overall error rate: 21.43%
```

```
set.seed(123)
forest2 <- rfsrc(acquisition ~ acq_exp_sq + industry + revenue + employees,
                 data = train.df,
                 importance = TRUE,
                 ntree = 1000)

forest2
```

```
##           Sample size: 350
##       Frequency of class labels: 105, 245
##           Number of trees: 1000
##       Forest terminal node size: 1
##       Average no. of terminal nodes: 53.859
## No. of variables tried at each split: 2
##       Total no. of variables: 4
##       Resampling used to grow trees: swor
##       Resample size used to grow trees: 221
##           Analysis: RF-C
##           Family: class
##       Splitting rule: gini *random*
##       Number of random split points: 10
##       Normalized brier score: 55.68
##           AUC: 85.55
##       Error rate: 0.18, 0.41, 0.09
##
## Confusion matrix:
##
##       predicted
## observed  0   1 class.error
##       0 62  43      0.4095
##       1 22 223      0.0898
##
## Overall error rate: 18.29%
```

```
set.seed(123)
forest3 <- rfsrc(acquisition ~ acq_exp + industry + revenue + employees,
                 data = train.df,
                 importance = TRUE,
                 ntree = 1000)

forest3
```

```
##              Sample size: 350
##      Frequency of class labels: 105, 245
##              Number of trees: 1000
##      Forest terminal node size: 1
##      Average no. of terminal nodes: 53.859
## No. of variables tried at each split: 2
##      Total no. of variables: 4
##      Resampling used to grow trees: swor
##      Resample size used to grow trees: 221
##              Analysis: RF-C
##              Family: class
##      Splitting rule: gini *random*
##      Number of random split points: 10
##      Normalized brier score: 55.68
##              AUC: 85.55
##      Error rate: 0.18, 0.41, 0.09
##
## Confusion matrix:
##
##      predicted
## observed  0  1 class.error
##      0 62  43      0.4095
##      1 22 223      0.0898
##
## Overall error rate: 18.29%
```

```
forest3$importance
```

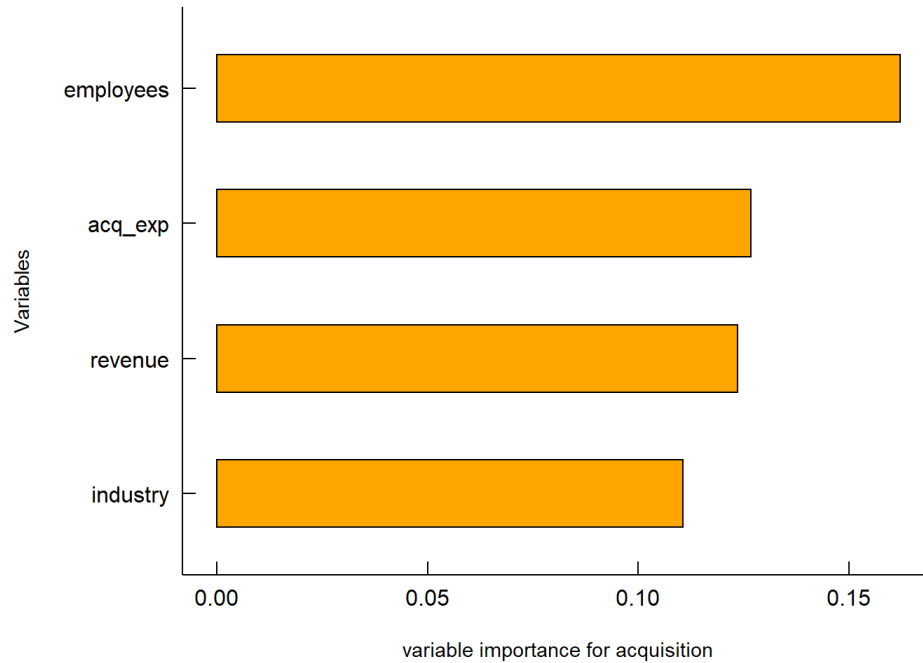
```
##              all              0              1
## acq_exp  0.02663842 0.27415814 -0.015089238
## industry 0.01059101 0.06006108  0.014978288
## revenue  0.02355819 0.19235080  0.008210321
## employees 0.06210708 0.55219954  0.002108055
```

```
forest3$importance[,1]
```

```
##      acq_exp  industry  revenue  employees
## 0.02663842 0.01059101 0.02355819 0.06210708
```

```
data.frame(importance = forest3$importance[,1] +.1) %>% # add a large +ve constant


tibble::rownames_to_column(var = "variable") %>%
ggplot(aes(x = reorder(variable,importance), y = importance)) +
  geom_bar(stat = "identity", fill = "orange", color = "black", width = 0.5)+
  coord_flip() +
  labs(x = "Variables", y = "variable importance for acquisition") +
  theme_nice
```



```
forest3$importance[,2]
```

```
##   acq_exp  industry  revenue  employees  
## 0.27415814 0.06006108 0.19235080 0.55219954
```


```
data.frame(importance = forest3$importance[,2] +.1) %>% # add a large +ve constant  
  
tibble::rownames_to_column(var = "variable") %>%  
ggplot(aes(x = reorder(variable,importance), y = importance)) +  
  geom_bar(stat = "identity", fill = "orange", color = "black", width = 0.5)+  
  coord_flip() +  
  labs(x = "Variables", y = "variable importance for acquisition") +  
  theme_nice
```



```
forest3$importance[,3]
```

```
##      acq_exp      industry      revenue      employees  
## -0.015089238  0.014978288  0.008210321  0.002108055
```

```
data.frame(importance = forest3$importance[,3] +.1) %>% # add a large +ve constant  
  
tibble::rownames_to_column(var = "variable") %>%  
ggplot(aes(x = reorder(variable, importance), y = importance)) +  
  geom_bar(stat = "identity", fill = "orange", color = "black", width = 0.5)+  
  coord_flip() +  
  labs(x = "Variables", y = "variable importance for acquisition") +  
  theme_nice
```



Optimized Selection

```
# Establish a List of possible values for hyper-parameters
mtry.values <- seq(1,4,1)
nodesize.values <- seq(1,4,1)
ntree.values <- seq(1e3,6e3,1e3)

# Create a data frame containing all combinations
hyper_grid <- expand.grid(mtry = mtry.values, nodesize = nodesize.values, ntree = ntree.values)

# Create an empty vector to store OOB error values
oob_err <- c()

# Write a Loop over the rows of hyper_grid to train the grid of models
for (i in 1:nrow(hyper_grid)) {
  # Train a Random Forest model
  set.seed(100)
  model <- rfsrc(acquisition ~ acq_exp_sq + industry + revenue + employees,
                 data = train.df,
                 mtry = hyper_grid$mtry[i],
                 nodesize = hyper_grid$nodesize[i],
                 ntree = hyper_grid$ntree[i])

  # Store OOB error for the model
  oob_err[i] <- model$err.rate[length(model$err.rate)]
}

# Identify optimal set of hyperparameters based on OOB error
opt_i <- which.min(oob_err)
print(hyper_grid[opt_i,])
```

```
## mtry nodesize ntree
## 1 1 1 1000
```

```
set.seed(111)
forest.hyper <- rfsrc(acquisition ~ acq_exp_sq + industry + revenue + employees,
                      data = train.df,
                      mtry = 1,
                      nodesize = 1,
                      ntree = 1000,
                      importance = TRUE)

forest.hyper
```

```
##              Sample size: 350
##      Frequency of class labels: 105, 245
##              Number of trees: 1000
##      Forest terminal node size: 1
##      Average no. of terminal nodes: 33.097
## No. of variables tried at each split: 1
##              Total no. of variables: 4
##      Resampling used to grow trees: swor
##      Resample size used to grow trees: 221
##              Analysis: RF-C
##              Family: class
##      Splitting rule: gini *random*
##      Number of random split points: 10
##      Normalized brier score: 54.08
##              AUC: 87.41
##      Error rate: 0.18, 0.45, 0.06
##
## Confusion matrix:
##
##      predicted
## observed  0   1 class.error
##      0 58  47      0.4476
##      1 15 230      0.0612
##
## Overall error rate: 17.71%
```

```
forest.hyper$importance
```


```
##              all              0              1
## acq_exp_sq 0.02751073 0.21487371 0.01364688
## industry   0.01357054 0.08620837 0.01520019
## revenue    0.02810808 0.20218839 0.02141340
## employees   0.06665774 0.49032627 0.04615532
```

```
forest.hyper$importance[,1]
```

```
## acq_exp_sq industry revenue employees
## 0.02751073 0.01357054 0.02810808 0.06665774
```

```
data.frame(importance = forest.hyper$importance[,1] +.1) %>% # add a large +ve constant

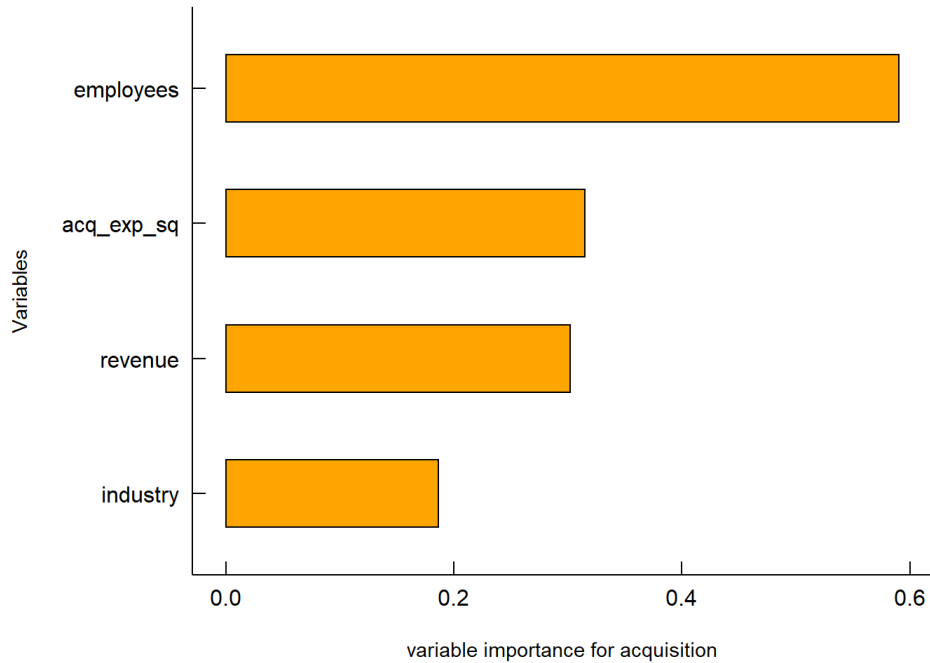
tibble::rownames_to_column(var = "variable") %>%
ggplot(aes(x = reorder(variable,importance), y = importance)) +
  geom_bar(stat = "identity", fill = "orange", color = "black", width = 0.5)+
  coord_flip() +
  labs(x = "Variables", y = "variable importance for acquisition") +
  theme_nice
```



```
forest.hyper$importance[,2]
```

```
## acq_exp_sq industry revenue employees  
## 0.21487371 0.08620837 0.20218839 0.49032627
```


```
data.frame(importance = forest.hyper$importance[,2] + .1) %>% # add a large +ve constant  
  
tibble::rownames_to_column(var = "variable") %>%  
ggplot(aes(x = reorder(variable, importance), y = importance)) +  
  geom_bar(stat = "identity", fill = "orange", color = "black", width = 0.5) +  
  coord_flip() +  
  labs(x = "Variables", y = "variable importance for acquisition") +  
  theme_nice
```

```
forest.hyper$importance[,3]
```

```
## acq_exp_sq industry revenue employees  
## 0.01364688 0.01520019 0.02141340 0.04615532
```

```
data.frame(importance = forest.hyper$importance[,3] + .1) %>% # add a large +ve constant  
  
tibble::rownames_to_column(var = "variable") %>%  
ggplot(aes(x = reorder(variable, importance), y = importance)) +  
  geom_bar(stat = "identity", fill = "orange", color = "black", width = 0.5) +  
  coord_flip() +  
  labs(x = "Variables", y = "variable importance for acquisition") +  
  theme_nice
```



```
error.df <- data.frame(pred1 = predict.rfsrc(forest3,newdata = test.df)$class,
  pred2 = predict.rfsrc(forest.hyper, newdata = test.df)$class,
  actual = test.df$acquisition)
```

```
PredsAll = predict.rfsrc(forest.hyper,newdata = acquisitionRetention)$class
```

```
Acquisition2.df <- cbind(acquisitionRetention,PredsAll)
```

```
Acquired.df <- filter(Acquisition2.df, PredsAll == "1")
```

Split for Duration prediction

```
set.seed(123)
idx.train_1 <- sample(1:nrow(Acquired.df), size = 0.7 * nrow(Acquired.df))
acq_train.df <- Acquired.df[idx.train_1,]
acq_test.df <- Acquired.df[-idx.train_1,]
```

```
mean(Acquired.df$duration)
```

```
## [1] 949.3203
```

Forest Duration Model

Manual Selection

```
set.seed(123)
forest_duration <- rfsrc(duration ~ profit + acq_exp + acq_exp_sq + ret_exp + ret_exp_sq + freq + freq_sq + crossbuy
  + sow + industry + revenue +employees,
  data = acq_train.df,
  importance = TRUE,
  ntree = 1000)

forest_duration
```

```
##           Sample size: 268
##           Number of trees: 1000
##           Forest terminal node size: 5
##           Average no. of terminal nodes: 28.273
## No. of variables tried at each split: 4
##           Total no. of variables: 12
##           Resampling used to grow trees: swor
##           Resample size used to grow trees: 169
##           Analysis: RF-R
##           Family: regr
##           Splitting rule: mse *random*
##           Number of random split points: 10
##           % variance explained: 99.29
##           Error rate: 1404.75
```


```
forest_duration$importance
```

```
##      profit      acq_exp  acq_exp_sq      ret_exp  ret_exp_sq      freq
## 1.202854e+04 8.132764e+01 5.792366e+01 2.829037e+04 2.864301e+04 1.949549e+04
##      freq_sq      crossbuy      sow      industry      revenue      employees
## 2.430503e+04 3.028284e+03 1.118423e+04 4.775466e-01 2.995912e+01 3.540989e+01
```


```
forest_duration$importance %>% log() # Log transform
```

```
##      profit      acq_exp  acq_exp_sq      ret_exp  ret_exp_sq      freq      freq_sq
## 9.3950371 4.3984859 4.0591259 10.2502768 10.2626649 9.8779386 10.0984388
##      crossbuy      sow      industry      revenue      employees
## 8.0157515 9.3222596 -0.7390936 3.3998339 3.5669911
```

```
data.frame(importance = forest_duration$importance + 100) %>% # add a large +ve constant
  log() %>%
  tibble::rownames_to_column(var = "variable") %>%
  ggplot(aes(x = reorder(variable, importance), y = importance)) +
    geom_bar(stat = "identity", fill = "orange", color = "black", width = 0.5) +
    coord_flip() +
    labs(x = "Variables", y = "Log-transformed variable importance for Duration") +
    theme_nice
```



```
data.frame(err.rate = forest_duration$err.rate) %>%
  na.omit() %>%
  tibble::rownames_to_column(var = "trees") %>%
  mutate(trees = as.numeric(trees)) %>%
  ggplot(aes(x = trees, y = err.rate, group = 1))+
  geom_line()+
  scale_x_continuous(breaks = seq(0,1250,100))+
  labs(x = "Number of trees", y = "OOB Error rate")+
  theme_nice
```



Optimized Selection

```

# Establish a list of possible values for hyper-parameters
mtry.values <- seq(1,12,1)
nodesize.values <- seq(1,5,1)
ntree.values <- seq(1e3,6e3,1e3)

# Create a data frame containing all combinations
hyper_grid <- expand.grid(mtry = mtry.values, nodesize = nodesize.values, ntree = ntree.values)

# Create an empty vector to store OOB error values
oob_err <- c()

# Write a loop over the rows of hyper_grid to train the grid of models
for (i in 1:nrow(hyper_grid)) {
  # Train a Random Forest model
  set.seed(123)
  model <- rfsrc(duration ~ profit + acq_exp + acq_exp_sq + ret_exp + ret_exp_sq + freq + freq_sq + crossbuy + sow +
industry + revenue + employees,
                data = acq_train.df,
                mtry = hyper_grid$mtry[i],
                nodesize = hyper_grid$nodesize[i],
                ntree = hyper_grid$ntree[i])

  # Store OOB error for the model
  oob_err[i] <- model$err.rate[length(model$err.rate)]
}

# Identify optimal set of hyperparameters based on OOB error
opt_i <- which.min(oob_err)
print(hyper_grid[opt_i,])

```

```

##      mtry nodesize ntree
## 6         6         1 1000

```

```

set.seed(100)

forest.hyper_duration <- rfsrc(duration ~ profit + acq_exp + acq_exp_sq + ret_exp + ret_exp_sq + freq + freq_sq + cro
ssbuy + sow + industry + revenue + employees,
                              data = acq_train.df,
                              mtry = 6,
                              nodesize = 1,
                              ntree = 1000,
                              importance = TRUE)

forest.hyper_duration

```

```

##              Sample size: 268
##              Number of trees: 1000
##              Forest terminal node size: 1
##              Average no. of terminal nodes: 140.294
## No. of variables tried at each split: 6
##              Total no. of variables: 12
##              Resampling used to grow trees: swor
##              Resample size used to grow trees: 169
##              Analysis: RF-R
##              Family: regr
##              Splitting rule: mse *random*
##              Number of random split points: 10
##              % variance explained: 99.47
##              Error rate: 1033.72

```

Interaction

```
find.interaction(forest.hyper_duration,  
                 method = "vimp",  
                 importance = "permute")
```

```
## Pairing ret_exp with ret_exp_sq
## Pairing ret_exp with freq
## Pairing ret_exp with freq_sq
## Pairing ret_exp with sow
## Pairing ret_exp with profit
## Pairing ret_exp with crossbuy
## Pairing ret_exp with acq_exp
## Pairing ret_exp with employees
## Pairing ret_exp with acq_exp_sq
## Pairing ret_exp with revenue
## Pairing ret_exp with industry
## Pairing ret_exp_sq with freq
## Pairing ret_exp_sq with freq_sq
## Pairing ret_exp_sq with sow
## Pairing ret_exp_sq with profit
## Pairing ret_exp_sq with crossbuy
## Pairing ret_exp_sq with acq_exp
## Pairing ret_exp_sq with employees
## Pairing ret_exp_sq with acq_exp_sq
## Pairing ret_exp_sq with revenue
## Pairing ret_exp_sq with industry
## Pairing freq with freq_sq
## Pairing freq with sow
## Pairing freq with profit
## Pairing freq with crossbuy
## Pairing freq with acq_exp
## Pairing freq with employees
## Pairing freq with acq_exp_sq
## Pairing freq with revenue
## Pairing freq with industry
## Pairing freq_sq with sow
## Pairing freq_sq with profit
## Pairing freq_sq with crossbuy
## Pairing freq_sq with acq_exp
## Pairing freq_sq with employees
## Pairing freq_sq with acq_exp_sq
## Pairing freq_sq with revenue
## Pairing freq_sq with industry
## Pairing sow with profit
## Pairing sow with crossbuy
## Pairing sow with acq_exp
## Pairing sow with employees
## Pairing sow with acq_exp_sq
## Pairing sow with revenue
## Pairing sow with industry
## Pairing profit with crossbuy
## Pairing profit with acq_exp
## Pairing profit with employees
## Pairing profit with acq_exp_sq
## Pairing profit with revenue
## Pairing profit with industry
## Pairing crossbuy with acq_exp
## Pairing crossbuy with employees
## Pairing crossbuy with acq_exp_sq
## Pairing crossbuy with revenue
## Pairing crossbuy with industry
## Pairing acq_exp with employees
## Pairing acq_exp with acq_exp_sq
## Pairing acq_exp with revenue
## Pairing acq_exp with industry
## Pairing employees with acq_exp_sq
## Pairing employees with revenue
## Pairing employees with industry
## Pairing acq_exp_sq with revenue
## Pairing acq_exp_sq with industry
## Pairing revenue with industry
##
```

```

##                               Method: vimp
##                               No. of variables: 12
##                               Variables sorted by VIMP?: TRUE
##                               No. of variables used for pairing: 12
##                               Total no. of paired interactions: 66
##                               Monte Carlo replications: 1
##                               Type of noising up used for VIMP: permute
##
##                               Var 1      Var 2      Paired      Additive      Difference
## ret_exp:ret_exp_sq      29162.0674  27781.6222  77061.7334  56943.6895  20118.0438
## ret_exp:freq      29162.0674  24575.8698  68111.5295  53737.9372  14373.5923
## ret_exp:freq_sq      29162.0674  21230.0776  61562.8800  50392.1449  11170.7350
## ret_exp:sow      29162.0674  9841.0873  44130.0836  39003.1547  5126.9290
## ret_exp:profit      29162.0674  7610.6515  44777.7263  36772.7189  8005.0075
## ret_exp:crossbuy      29162.0674  2439.9228  32714.6475  31601.9902  1112.6573
## ret_exp:acq_exp      29162.0674  8.7771  29265.8887  29170.8444  95.0442
## ret_exp:employees      29162.0674  7.0004  29008.1172  29169.0677  -160.9505
## ret_exp:acq_exp_sq      29162.0674  14.1276  29386.9569  29176.1950  210.7619
## ret_exp:revenue      29162.0674  -10.1128  29706.2562  29151.9546  554.3016
## ret_exp:industry      29162.0674  -0.3129  29377.5861  29161.7545  215.8317
## ret_exp_sq:freq      27689.5693  25140.5006  64703.7923  52830.0699  11873.7223
## ret_exp_sq:freq_sq      27689.5693  21331.3289  58757.0650  49020.8982  9736.1668
## ret_exp_sq:sow      27689.5693  10049.1674  43547.1938  37738.7367  5808.4571
## ret_exp_sq:profit      27689.5693  7587.2712  43251.3910  35276.8405  7974.5505
## ret_exp_sq:crossbuy      27689.5693  2289.6482  31549.0636  29979.2175  1569.8462
## ret_exp_sq:acq_exp      27689.5693  7.7812  28074.2247  27697.3505  376.8742
## ret_exp_sq:employees      27689.5693  24.7080  27577.2162  27714.2773  -137.0611
## ret_exp_sq:acq_exp_sq      27689.5693  10.1858  27051.9374  27699.7551  -647.8177
## ret_exp_sq:revenue      27689.5693  -25.2575  27628.5857  27664.3118  -35.7260
## ret_exp_sq:industry      27689.5693  13.0283  27532.6265  27702.5976  -169.9710
## freq:freq_sq      24329.7513  21244.3025  53648.7598  45574.0538  8074.7059
## freq:sow      24329.7513  9659.9582  36963.4081  33989.7095  2973.6986
## freq:profit      24329.7513  7475.8841  35696.1836  31805.6354  3890.5482
## freq:crossbuy      24329.7513  2385.8940  28016.3055  26715.6453  1300.6603
## freq:acq_exp      24329.7513  21.9528  24392.1532  24351.7041  40.4490
## freq:employees      24329.7513  17.3702  23649.8764  24347.1215  -697.2451
## freq:acq_exp_sq      24329.7513  11.7574  23659.1821  24341.5087  -682.3266
## freq:revenue      24329.7513  -11.2113  23886.3653  24318.5400  -432.1747
## freq:industry      24329.7513  3.2423  24469.8930  24332.9936  136.8994
## freq_sq:sow      20928.6061  9381.4070  33871.8585  30310.0131  3561.8454
## freq_sq:profit      20928.6061  7410.6120  33001.0761  28339.2181  4661.8580
## freq_sq:crossbuy      20928.6061  2222.7622  24322.7717  23151.3683  1171.4034
## freq_sq:acq_exp      20928.6061  20.9258  21075.5170  20949.5319  125.9851
## freq_sq:employees      20928.6061  9.4222  21214.3623  20938.0283  276.3340
## freq_sq:acq_exp_sq      20928.6061  15.5764  20675.1022  20944.1825  -269.0803
## freq_sq:revenue      20928.6061  -6.5836  20607.3730  20922.0225  -314.6495
## freq_sq:industry      20928.6061  7.8855  20151.4441  20936.4916  -785.0475
## sow:profit      9547.3327  7809.2866  19778.2804  17356.6193  2421.6611
## sow:crossbuy      9547.3327  2283.3808  12485.3739  11830.7135  654.6604
## sow:acq_exp      9547.3327  4.1886  9631.0344  9551.5213  79.5131
## sow:employees      9547.3327  6.7578  9664.0344  9554.0905  109.9439
## sow:acq_exp_sq      9547.3327  32.3528  9613.5027  9579.6855  33.8172
## sow:revenue      9547.3327  5.6753  9566.5262  9553.0080  13.5182
## sow:industry      9547.3327  9.8987  9330.2969  9557.2314  -226.9345
## profit:crossbuy      7731.5555  2370.0466  10382.3055  10101.6021  280.7034
## profit:acq_exp      7731.5555  15.6825  8011.9140  7747.2380  264.6760
## profit:employees      7731.5555  8.8034  7702.8616  7740.3589  -37.4973
## profit:acq_exp_sq      7731.5555  4.7379  8093.6608  7736.2934  357.3674
## profit:revenue      7731.5555  -17.2039  7823.2890  7714.3516  108.9374
## profit:industry      7731.5555  0.8602  7924.4565  7732.4157  192.0407
## crossbuy:acq_exp      2536.2765  20.0742  2643.6470  2556.3507  87.2963
## crossbuy:employees      2536.2765  19.7301  2451.3131  2556.0066  -104.6935
## crossbuy:acq_exp_sq      2536.2765  18.4907  2587.7264  2554.7672  32.9591
## crossbuy:revenue      2536.2765  -12.4676  2684.9829  2523.8089  161.1740
## crossbuy:industry      2536.2765  9.8129  2465.3070  2546.0894  -80.7824
## acq_exp:employees      19.3330  15.0031  36.5465  34.3361  2.2104
## acq_exp:acq_exp_sq      19.3330  38.0846  41.6116  57.4176  -15.8059
## acq_exp:revenue      19.3330  -14.2932  -13.7124  5.0398  -18.7521

```


## acq_exp:industry	19.3330	12.3581	52.6970	31.6911	21.0060
## employees:acq_exp_sq	11.0335	9.5527	29.5500	20.5862	8.9639
## employees:revenue	11.0335	-3.4563	24.1585	7.5772	16.5814
## employees:industry	11.0335	10.8080	18.9534	21.8415	-2.8881
## acq_exp_sq:revenue	41.8106	-10.2718	40.2954	31.5389	8.7565
## acq_exp_sq:industry	41.8106	11.5000	42.5628	53.3107	-10.7479
## revenue:industry	-3.7107	12.1593	24.5734	8.4485	16.1249

```
# Establish a list of possible values for hyper-parameters
mtry.values <- seq(1,12,1)
nodesize.values <- seq(1,5,1)
ntree.values <- seq(1e3,6e3,1e3)

# Create a data frame containing all combinations
hyper_grid <- expand.grid(mtry = mtry.values, nodesize = nodesize.values, ntree = ntree.values)

# Create an empty vector to store OOB error values
oob_err <- c()

# Write a loop over the rows of hyper_grid to train the grid of modelsfor (i in 1:nrow(hyper_grid)) {
  for (i in 1:nrow(hyper_grid)) {
    # Train a Random Forest model
    set.seed(123)
    model <- rfsrc(duration ~ profit + acq_exp + acq_exp_sq + ret_exp + ret_exp_sq + freq + freq_sq + crossbuy + sow +
industry + revenue +employees + ret_exp*ret_exp_sq + ret_exp*freq + ret_exp*freq_sq + ret_exp_sq*freq + ret_exp_sq*fr
eq_sq + freq*freq_sq ,
                    data = acq_train.df,
                    mtry = hyper_grid$mtry[i],
                    nodesize = hyper_grid$nodesize[i],
                    ntree = hyper_grid$ntree[i])

    # Store OOB error for the model
    oob_err[i] <- model$err.rate[length(model$err.rate)]
  }

# Identify optimal set of hyperparameters based on OOB error
opt_i <- which.min(oob_err)
print(hyper_grid[opt_i,])
```

```
## mtry nodesize ntree
## 6 6 1 1000
```

```
set.seed(100)

forest.hyper_duration2 <- rfsrc(duration ~ profit + acq_exp + acq_exp_sq + ret_exp + ret_exp_sq + freq + freq_sq + c
rossbuy + sow + industry + revenue +employees + ret_exp*ret_exp_sq + ret_exp*freq + ret_exp*freq_sq + ret_exp_sq*freq
+ ret_exp_sq*freq_sq + freq*freq_sq,
                                data = acq_train.df,
                                mtry = 6,
                                nodesize = 1,
                                ntree = 1000,
                                importance = TRUE)

forest.hyper_duration2
```

```
##           Sample size: 268
##           Number of trees: 1000
##           Forest terminal node size: 1
##           Average no. of terminal nodes: 140.294
## No. of variables tried at each split: 6
##           Total no. of variables: 12
##           Resampling used to grow trees: swor
##           Resample size used to grow trees: 169
##           Analysis: RF-R
##           Family: regr
##           Splitting rule: mse *random*
##           Number of random split points: 10
##           % variance explained: 99.47
##           Error rate: 1033.72
```

```
PredsDuration = predict.rfsrc(forest.hyper_duration,newdata = acq_test.df)$predicted
DurationDF <- data.frame(acq_test.df$duration, PredsDuration)
mse(acq_test.df$duration, PredsDuration)
```

```
## [1] 1525.854
```

```
MAE_D<-MAE(acq_test.df$duration, PredsDuration)
MAE_D
```

```
## [1] 25.36803
```

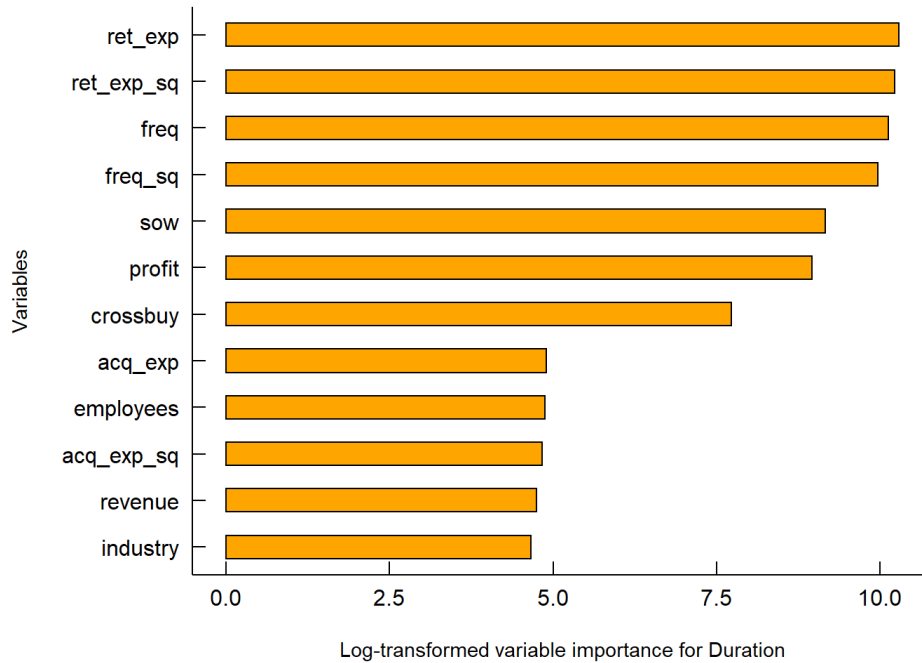
```
forest.hyper_duration$importance
```

```
##      profit      acq_exp  acq_exp_sq      ret_exp  ret_exp_sq      freq
## 7654.818215  33.448018   25.895889 29288.258733 27470.362939 24892.665145
##      freq_sq  crossbuy      sow  industry  revenue  employees
## 21207.728904 2166.546657 9480.508392  6.043773  15.641571  30.954803
```


```
forest.hyper_duration$importance %>% log() # Log transform
```

```
##      profit      acq_exp  acq_exp_sq      ret_exp  ret_exp_sq      freq  freq_sq
##  8.943091  3.509993  3.254084  10.284942  10.220863  10.122328  9.962121
##  crossbuy      sow  industry  revenue  employees
##  7.680890  9.156993  1.799028  2.749932  3.432528
```

```
data.frame(importance = forest.hyper_duration$importance + 100) %>% # add a Large +ve constant
log() %>%
tibble::rownames_to_column(var = "variable") %>%
ggplot(aes(x = reorder(variable,importance), y = importance)) +
  geom_bar(stat = "identity", fill = "orange", color = "black", width = 0.5)+
  coord_flip() +
  labs(x = "Variables", y = "Log-transformed variable importance for Duration") +
  theme_nice
```



```
data.frame(err.rate = forest.hyper_duration$err.rate) %>%
  na.omit() %>%
  tibble::rownames_to_column(var = "trees") %>%
  mutate(trees = as.numeric(trees)) %>%
  ggplot(aes(x = trees, y = err.rate, group = 1))+
  geom_line()+
  scale_x_continuous(breaks = seq(0,1250,100))+
  labs(x = "Number of trees", y = "OOB Error rate")+
  theme_nice
```



PDP Plots

Duration: Retention Expenditure

```
min(forest.hyper_duration$xvar$ret_exp)
```

```
## [1] 0
```

```
max(forest.hyper_duration$xvar$ret_exp)
```

```
## [1] 1082.37
```


```
ret_exp_seq = seq(0,1100,20)
```

```
# extract marginal effect using partial dependence
marginal.effect <- partial(forest.hyper_duration,
  partial.xvar = "ret_exp",
  partial.values = ret_exp_seq)

means.exp <- marginal.effect$regrOutput$duration %>% colMeans()
```

```
marginal.effect.df <-
  data.frame(pred.duration = means.exp, ret_exp_seq = ret_exp_seq)
```


```
ggplot(marginal.effect.df, aes(x = ret_exp_seq, y = pred.duration)) +
  geom_point(shape = 21, color = "purple", size = 2, stroke = 1.2)+
  geom_smooth(method = "lm", formula = y ~ poly(x,3), se = FALSE, color = "black")+ # try with other values
  labs(x = "Average retention in $", y = "Predicted duration") +
  scale_x_continuous(breaks = seq(0,1100,20))+
  theme_nice # positive effect of ret_exp not clear as suggested by reg coefs
```



```
# first check relationship between actual duration and ret_exp
```

```
ggplot(acquisitionRetention, aes(x = ret_exp, y = duration)) +
  geom_point(shape = 21, col = "purple", size = 3) +
  stat_smooth(method = "lm", se = FALSE, color = "black") +
  scale_x_continuous(breaks = seq(0,1100,100)) +
  scale_y_continuous(breaks = seq(0,2200,200)) +
  geom_rug(sides = "b", col = "red", alpha = 0.2) +
  labs(y = "Actual duration", x = "retention in $") +
  theme_nice
```

```
## `geom_smooth()` using formula 'y ~ x'
```



```
# repeat with smaller values of ret_exp
```


```
ret_exp_seq2 = seq(200,700,25)
```

```
marginal.effect.new <- partial(forest.hyper_duration,
  partial.xvar = "ret_exp",
  partial.values = ret_exp_seq2)
```

```
means.exp.new <- marginal.effect.new$regrOutput$duration %>% colMeans()
```


```
marginal.effect.df.new <-
  data.frame(pred.duration = means.exp.new, ret_exp_seq = ret_exp_seq2)
```

```
ggplot(marginal.effect.df.new, aes(x = ret_exp_seq, y = pred.duration)) +
  geom_point(shape = 21, color = "purple", size = 2, stroke = 1.2)+
  geom_path()+
  labs(x = "retention in $", y = "Predicted duration") +
  scale_x_continuous(breaks = seq(200,700,25))+
  theme_nice
```



Duration: Crossbuy Categories

```
CrossbuyPlot <- boxplot(data=acquisitionRetention, duration ~ crossbuy)
```



```
x <- c("Minimum", "25th Percentile", "Median", "75th Percentile", "Maximum")
```

```
CrossbuyStats <- data.frame(x, CrossbuyPlot$stats)
```

```
colnames(CrossbuyStats) <- c("Statistic","0","1","2","3","4","5","6","7", "8", "9", "10", "11")
```

Duration: Retention Expenditure Squared

```
min(forest.hyper_duration$xvar$ret_exp_sq)
```

```
## [1] 0
```

```
max(forest.hyper_duration$xvar$ret_exp_sq)
```

```
## [1] 1171525
```


```
ret_sq_seq = seq(0,1200000,100000)

# extract marginal effect using partial dependence
marginal.effect <- partial(forest.hyper_duration,
                           partial.xvar = "ret_exp_sq",
                           partial.values = ret_sq_seq)

means.exp <- marginal.effect$regrOutput$duration %>% colMeans()

marginal.effect.df <-
  data.frame(pred.duration = means.exp, ret_sq_seq = ret_sq_seq)
```


```
ggplot(marginal.effect.df, aes(x = ret_sq_seq, y = pred.duration)) +
  geom_point(shape = 21, color = "purple", size = 2, stroke = 1.2)+
  geom_smooth(method = "lm", formula = y ~ poly(x,3), se = FALSE, color = "black")+ # try with other values
  labs(x = "Squared Retention Expenditure", y = "Predicted duration") +
  scale_x_continuous(breaks = seq(0,1200000,100000))+
  theme_nice # positive effect of ret_exp not clear as suggested by reg coefs
```



```
# first check relationship between actual duration and ret_exp
```

```
ggplot(acquisitionRetention, aes(x = ret_exp_sq, y = duration)) +
  geom_point(shape = 21, col = "purple", size = 3) +
  stat_smooth(method = "lm", se = FALSE, color = "black") +
  scale_x_continuous(breaks = seq(0,1200000,100000)) +
  scale_y_continuous(breaks = seq(0,1800,100)) +
  geom_rug(sides = "b", col = "red", alpha = 0.2) +
  labs(y = "Actual duration", x = "Squared Retention Expenditure") +
  theme_nice
```

```
## `geom_smooth()` using formula 'y ~ x'
```




```
# repeat with smaller values of ret_exp
ret_sq_seq2 = seq(50000,500000,25000)

marginal.effect.new <- partial(forest.hyper_duration,
                              partial.xvar = "ret_exp_sq",
                              partial.values = ret_sq_seq2)

means.exp.new <- marginal.effect.new$regrOutput$duration %>% colMeans()

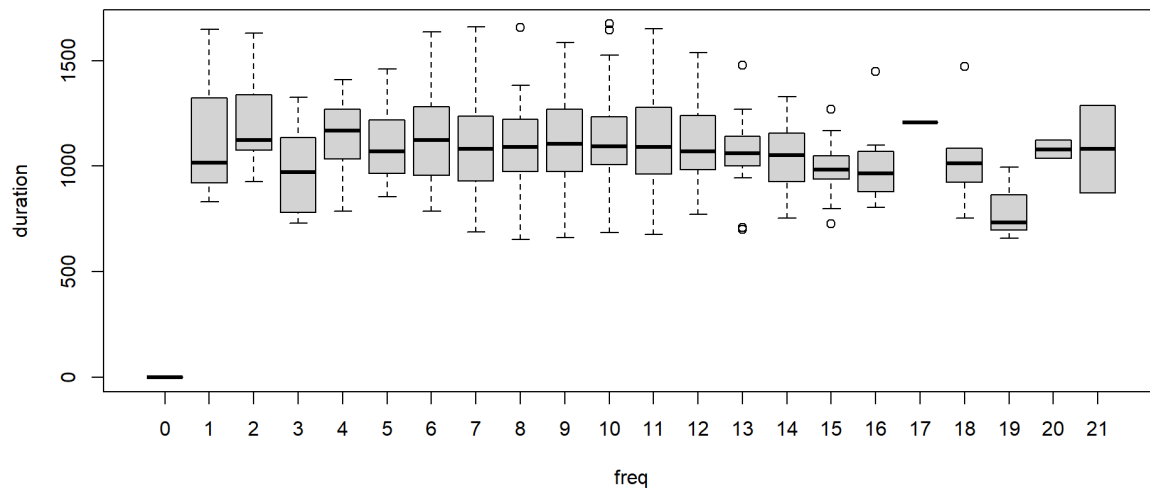
marginal.effect.df.new <-
  data.frame(pred.duration = means.exp.new, ret_sq_seq = ret_sq_seq2)

ggplot(marginal.effect.df.new, aes(x = ret_sq_seq, y = pred.duration)) +
  geom_point(shape = 21, color = "purple", size = 2, stroke = 1.2)+
  geom_path()+
  labs(x = "Squard Retention Expenditure", y = "Predicted duration") +
  scale_x_continuous(breaks = seq(50000,500000,25000))+
  theme_nice
```



Duration: frequency

```
FreqPlot <- boxplot(data=acquisitionRetention, duration ~ freq)
```




```
x <- c("Minimum", "25th Percentile", "Median", "75th Percentile", "Maximum")
```

```
FreqStats <- data.frame(x, FreqPlot$stats)
```

```
colnames(FreqStats) <- c("Statistic", "0", "1", "2", "3", "4", "5", "6", "7", "8", "9", "10", "11", "12", "13", "14", "15",  
"16", "17", "18", "19", "20", "21")
```

Duration: frequency squared

```
FreqSqPlot <- boxplot(data=acquisitionRetention, duration ~ freq_sq)
```



```
x <- c("Minimum", "25th Percentile", "Median", "75th Percentile", "Maximum")

FreqSqStats <- data.frame(x, FreqSqPlot$stats)
colnames(FreqSqStats) <- c("Statistic", "0", "1", "4", "9", "16", "25", "36", "49", "64", "81", "100", "121", "144", "169",
"196", "225", "256", "289", "324", "361", "400", "441")
```


Minimal Depth

```
mindepth <- max.subtree(forest.hyper_duration, sub.order = TRUE)

print(round(mindepth$order, 3)[,1])
```

```
##      profit    acq_exp acq_exp_sq    ret_exp ret_exp_sq      freq  freq_sq
##      2.854      4.793      4.860      1.544      1.562      2.728      2.865
##  crossbuy      sow    industry    revenue  employees
##      4.478      3.974      5.911      4.736      4.659
```

```
data.frame(md = round(mindepth$order, 3)[,1]) %>%
  tibble::rownames_to_column(var = "variable") %>%
  ggplot(aes(x = reorder(variable, desc(md)), y = md)) +
    geom_bar(stat = "identity", fill = "orange", color = "black", width = 0.2) +
    coord_flip() +
    labs(x = "Variables", y = "Minimal Depth") +
    theme_nice
```



```
as.matrix(mindepth$sub.order) %>%
  reshape2::melt() %>%
  data.frame() %>%
  ggplot(aes(x = Var1, y = Var2, fill = value)) +
    scale_x_discrete(position = "top") +
    geom_tile(color = "white") +
    viridis::scale_fill_viridis("Relative min. depth") +
    labs(x = "", y = "") +
    theme_bw()
```

